

MED-NEXUS

A-to-Z Full Research Explanation

Implementation • Methodology • Project 2 • Project 4 • Experiments • Results • Efficiency • Limitations • Journal Handover

Purpose	A complete speaking and writing guide for explaining the current Med-Nexus research implementation and documented experimental results to a journal/results writer or supervisor.
Primary source	Med-Nexus_Complete_Research_Documentation (current final research documentation).
Final evaluation	1× NVIDIA A100-SXM4 80GB • 5 seeds • N=5,000 held-out paired image-text clinical cases.

How to use this document: Read Sections 1–12 for the research story, Sections 13–18 for results and interpretation, and Sections 19–25 before the journal call. The final page contains a compact call script.

Contents

1. Executive summary
2. What problem Med-Nexus addresses
3. Project 2: retrieval-aware robustness
4. Project 4: intervention-based attention contribution
5. How Project 2 + Project 4 become Med-Nexus
6. End-to-end architecture
7. What each component does
8. Research questions and hypotheses
9. Experimental protocol
10. Main ablation: MN-RUN-01 to MN-RUN-06
11. Final diagnostic results
12. Grounding, retrieval, conflict and calibration
13. Δ Loss explained carefully
14. Efficiency and contribution-aware gating
15. What the experiments demonstrate
16. What the results do not establish
17. Reproducibility and repository handover
18. What to say about the dataset and evaluation set
19. Questions the journal writer may ask
20. Safe scientific wording
21. Five-minute A-to-Z call script
22. One-minute quick reference
23. Glossary
24. Final checklist before the call
25. Source-of-truth note

1. Executive Summary

Med-Nexus is the application-level medical multimodal research system described in the current final documentation. It integrates two foundational research directions: Project 2, which contributes retrieval-aware multimodal robustness, evidence handling and cross-modal conflict analysis; and Project 4, which contributes intervention-based attention contribution analysis and contribution-aware gating for efficient long-context inference.

The project is therefore not only a diagnostic classifier. Its evaluation is designed to examine four connected dimensions: diagnostic discrimination, evidence/retrieval behaviour, visual-textual conflict, and computational efficiency. Calibration metrics are included so that predictive confidence can also be assessed.

One-sentence explanation: Med-Nexus studies whether a medical multimodal system can use retrieved evidence more reliably, explicitly detect disagreement between modalities, measure which attention interactions contribute to the task, and use that contribution information to reduce unnecessary long-context computation.

The current final documentation reports a five-seed evaluation using one NVIDIA A100-SXM4 80GB on Vast.ai and a held-out paired image-text clinical evaluation set of N=5,000. The final Full Med-Nexus configuration is 33.40B parameters.

2. What Problem Med-Nexus Addresses

A conventional medical imaging system can be described as image → model → prediction. A multimodal system adds text or retrieved clinical evidence, but this introduces new failure modes: retrieved evidence can be noisy or irrelevant, visual inputs can be ambiguous, and image-derived information can disagree with retrieved text.

Med-Nexus treats these issues as measurable research questions rather than assuming that retrieval is always correct or that attention overlap automatically identifies redundant information.

- Evidence quality: Does the system remain reliable when retrieval is noisy?
- Visual ambiguity: How does performance change when visual evidence is degraded?
- Cross-modal conflict: Can disagreement between image-derived information and retrieved clinical evidence be identified?
- Attention contribution: Which interactions actually matter to the task?
- Efficiency: Can contribution information guide computation without relying on overlap alone?
- Calibration: Are confidence estimates better behaved as the system changes?

The final research documentation explicitly frames the work as research evaluation under the stated experimental conditions, not as clinical deployment evidence.

3. Project 2 — Retrieval-Aware Robustness

Project 2 is the reliability/evidence side of Med-Nexus. In the medical application, retrieval is treated as an information source whose quality and consistency can be challenged.

The project operationalizes Project 2 through retrieval noise, visual ambiguity, cross-modal conflict, evidence reliability and calibration.

Retrieval noise

The system evaluates behaviour when retrieved evidence is not perfectly clean. The scientific question is not simply whether retrieval improves an average score; it is whether the model behaves robustly when evidence quality is challenged.

Visual ambiguity

The system includes visual stress conditions such as ambiguity/degradation in the evaluation design. The purpose is to test how multimodal behaviour changes when the visual signal becomes less certain.

Cross-modal conflict

A conflict occurs when visual evidence and retrieved clinical evidence disagree. The system explicitly evaluates conflict detection rather than assuming the two sources are always consistent.

Simple example: Image-derived evidence suggests an abnormal finding, while retrieved text supports a normal finding. The important research question is whether the system can detect and handle that disagreement.

4. Project 4 — Intervention-Based Attention Contribution

Project 4 is the attention-efficiency side of Med-Nexus. Its central conceptual point is that attention overlap should not automatically be treated as redundancy.

Two interactions can look similar in attention space while having different task importance. Therefore, Med-Nexus evaluates contribution through intervention: mask or intervene on an interaction and measure the change in task loss.

Documented interaction quantity:

$$\Delta_{ij} = L(\text{mask}(i,j)) - L(\text{dense})$$

Here, the masked/intervened condition is compared with the dense reference. This creates a task-oriented contribution signal instead of treating attention overlap as a direct synonym for redundancy.

The contribution information is then connected to the long-context gating pathway, where the goal is to reduce unnecessary computation according to measured contribution.

5. How Project 2 + Project 4 Become Med-Nexus

The application-level logic is:

Foundational idea	Med-Nexus interpretation
Project 2	Retrieval-aware evidence handling and robustness under noisy/ambiguous/conflicting evidence.
Project 4	Intervention-based contribution measurement and contribution-aware long-context computation.
Med-Nexus	Medical multimodal evaluation combining evidence reliability, conflict analysis, diagnostic prediction, ca

The key research story: Project 2 answers “Can the system recognize when its evidence is unreliable or contradictory?” Project 4 answers “Which attention interactions actually contribute to the task, and can that information be used to reduce computation?” Med-Nexus evaluates both together in a medical multimodal setting.

6. End-to-End Architecture

Stage	Input / operation	Output
1	Medical image + clinical/query text	Multimodal input
2	Visual representation	Image features
3	Clinical evidence retrieval	Retrieved evidence
4	Evidence reliability / conflict analysis	Evidence/conflict signals
5	Multimodal fusion	Joint representation
6	Contribution-aware long-context attention	Contribution/gating signals
7	Diagnostic prediction	Prediction + confidence
8	Evaluation	Diagnostic + calibration + retrieval + conflict + efficiency metrics

In compact form:

Medical image + clinical/query text → visual representation → clinical evidence retrieval → evidence reliability/conflict analysis → multimodal fusion → contribution-aware long-context attention → diagnostic prediction → calibration/reliability evaluation.

7. What Each Component Does

Vision encoder: Converts the medical image into a learned visual representation. The implementation uses DenseNet-121.

Retrieval: Finds clinical evidence relevant to the multimodal input and provides the evidence stream used by the reliability/conflict experiments.

Text / clinical encoding: Represents retrieved or clinical text so it can participate in multimodal fusion.

Conflict detection: Measures disagreement between image-derived and retrieved evidence.

Multimodal fusion: Combines visual and textual/evidence representations.

Intervention analysis: Masks/intervenes on attention interactions and measures task-loss change.

Contribution-aware gating: Uses contribution information in the long-context computation pathway.

Prediction head: Produces the diagnostic output used for evaluation.

Calibration evaluation: Reports ECE and Brier score alongside discrimination metrics.

8. Research Questions and Hypotheses

RQ1	Retrieval-aware evidence handling	How does multimodal behaviour change when evidence quality changes?
RQ2	Cross-modal conflict detection	Can disagreement between visual and retrieved evidence be explicitly detected?
RQ3	Intervention-based contribution	Can task-loss change from attention intervention provide a useful contribution signal?
RQ4	Contribution-aware efficiency	Can contribution information support reduced long-context computation?
RQ5	Scale + mechanisms	How do model scale, retrieval, conflict handling and gating affect performance and calibration?

9. Experimental Protocol

Item	Final documented protocol
Hardware	1× NVIDIA A100-SXM4 80GB on Vast.ai
Seeds	42, 123, 456, 789, 2024
Held-out evaluation	N=5,000 paired image-text clinical cases
Maximum sequence length	2,048
Precision	AMP-bf16
Batch size	32/GPU
Effective batch	64
Training	100 epochs; 12,500 steps; 1,000 warmup steps
Optimizer	AdamW
β_1 / β_2	0.9 / 0.999
Weight decay	0.05
Peak learning rate	3×10^{-4}

These values are reproduced from the current complete research documentation. If the journal requires dataset provenance, patient-level split details, or exact run-log identifiers, those should be taken from the underlying experiment records rather than guessed.

10. Main Ablation — MN-RUN-01 to MN-RUN-06

Run	Configuration	Params	AUROC	AUPRC	F1	Accuracy
MN-RUN-01	8B Baseline	8.12B	0.812 ± 0.004	0.745 ± 0.006	0.738 ± 0.005	0.795 ± 0.004
MN-RUN-02	12B Baseline	12.35B	0.849 ± 0.003	0.789 ± 0.005	0.775 ± 0.004	0.824 ± 0.003
MN-RUN-03	32B Dense baseline	32.10B	0.881 ± 0.003	0.832 ± 0.004	0.814 ± 0.003	0.856 ± 0.003
MN-RUN-04	+ overlap suppression	32.65B	0.918 ± 0.002	0.882 ± 0.003	0.859 ± 0.002	0.889 ± 0.002
MN-RUN-05	+ intervention Δ Loss	33.12B	0.941 ± 0.002	0.912 ± 0.002	0.887 ± 0.002	0.915 ± 0.002
MN-RUN-06	Full Med-Nexus	33.40B	0.962 ± 0.001	0.941 ± 0.002	0.918 ± 0.001	0.938 ± 0.001

How to explain this table: The first three rows establish model-capacity baselines. The later rows progressively introduce the attention-related research mechanisms, ending in the full Med-Nexus configuration. The table is an ablation: its purpose is to attribute measured differences to controlled configuration changes.

11. Final Diagnostic Results

Metric	Full Med-Nexus
Parameters	33.40B
AUROC	0.962 ± 0.001
AUPRC	0.941 ± 0.002
F1	0.918 ± 0.001

Metric	Full Med-Nexus
Sensitivity	0.912 ± 0.002
Specificity	0.954 ± 0.001
Accuracy	0.938 ± 0.001

Interpretation: within the documented evaluation, the Full Med-Nexus configuration reports the final ablation values above. These should be described as measured experimental results under the declared protocol, not as universal performance guarantees.

AUROC measures discrimination across thresholds; AUPRC summarizes precision-recall behaviour; F1 combines precision and recall; sensitivity measures true-positive recovery; specificity measures true-negative recovery; accuracy is the overall fraction correctly classified.

12. Grounding, Retrieval, Conflict and Calibration

Run	ΔLoss	Retrieval Top-1 / Top-5	Conflict detection	ECE	Brier
MN-RUN-03	—	—	—	0.052 ± 0.002	0.104 ± 0.003
MN-RUN-04	-0.142 ± 0.008	88.4% / 96.1%	82.5%	0.038 ± 0.001	0.081 ± 0.002
MN-RUN-05	-0.285 ± 0.005	93.7% / 98.8%	92.1%	0.021 ± 0.001	0.062 ± 0.001
MN-RUN-06	-0.341 ± 0.004	96.8% / 99.6%	97.4%	0.011 ± 0.001	0.045 ± 0.001

ECE: Expected Calibration Error. Lower values indicate smaller average calibration discrepancy under the reported evaluation.

Brier: a probabilistic prediction error measure. Lower values indicate better probabilistic scoring under the reported evaluation.

The final Full Med-Nexus configuration reports retrieval Top-1/Top-5 of 96.8%/99.6%, conflict detection of 97.4%, ECE of 0.011 ± 0.001, and Brier score of 0.045 ± 0.001.

13. ΔLoss Explained Carefully

Definition: $\Delta_{ij} = L(\text{mask}(i,j)) - L(\text{dense})$

For a selected interaction between positions i and j , the intervention removes or masks the interaction and compares the resulting task loss with the dense reference.

Under the documented definition, a negative ΔLoss means the intervened condition has lower loss than the clean/dense reference for the stated comparison. It should not be interpreted in isolation as a universal measure of causality, safety or clinical importance.

Full Med-Nexus: -0.341 ± 0.004 .

When explaining this to a paper writer, always describe the exact intervention and reference condition. The sign only has meaning relative to that definition.

14. Efficiency and Contribution-Aware Gating

Run	Latency	Inference VRAM	Token gating	Throughput
MN-RUN-03	21.4 ms/sample	5.2 GB	0%	46.7/s

Run	Latency	Inference VRAM	Token gating	Throughput
MN-RUN-04	23.1 ms/sample	5.5 GB	0%	43.2/s
MN-RUN-05	24.8 ms/sample	5.7 GB	0%	40.3/s
MN-RUN-06	11.2 ms/sample	3.1 GB	46.2%	89.3/s

Relative to MN-RUN-05, the final documentation derives a 54.8% latency reduction and a 45.6% inference-VRAM reduction for MN-RUN-06.

The important phrase is contribution-aware gating: the system uses contribution information in the long-context pathway to gate computation. Do not describe 46.2% token gating as deleting 46.2% of the model; it is the reported gating statistic.

15. What the Experiments Demonstrate

- The scaling ablation evaluates how diagnostic performance changes with model capacity.
- Retrieval and conflict experiments evaluate evidence handling and visual-textual inconsistency.
- The intervention experiment evaluates attention-edge contribution using measured task-loss change.
- Contribution gating evaluates whether contribution information can reduce computation while retaining task performance in the stated evaluation.
- The full-system ablation evaluates the combined system under the declared protocol.
- Calibration metrics provide an additional view of prediction reliability beyond discrimination metrics.

16. What the Results Do NOT Establish

This section should be repeated in the journal discussion/limitations where appropriate.

- Prospective clinical validation.
- Regulatory approval or authorization.
- Clinical deployment readiness.
- Replacement of a radiologist or clinician.
- Universal generalization to all medical imaging populations or modalities.
- Independent proof that the method is state-of-the-art against every relevant published baseline.
- That the reported benchmark results automatically translate into patient outcomes.

The current documentation explicitly states that the results are research evaluation results under the specified experimental conditions.

17. Reproducibility and Repository Handover

The project repository is intended to contain the implementation, configurations, experiments, tests, result tables/figures, API, documentation and Vast.ai execution material.

Area	Purpose
train.py	Training entry point.
evaluate.py	Checkpoint/held-out evaluation entry point.

Area	Purpose
models/	Vision, text, fusion and attention-related model components.
services/	Retrieval, conflict detection and inference services.
experiments/	Baseline, retrieval-noise, visual ambiguity, conflict, combined and Project 4 experiments.
evaluation/	Metrics, reporting and Project 4 statistics.
tests/	Automated software tests.
results/	Result tables, figures and documented audit artifacts.
docs/	Research reports, reproducibility and experiment documentation.
VAST_AI_MED_NEXUS.ipynb	Vast.ai execution material.

18. What to Say About the Dataset and Evaluation Set

The current final documentation supports the statement that the evaluation used a held-out paired image-text clinical test/evaluation set of N=5,000. It also documents the five seeds and experimental protocol.

For any journal question asking for the exact public dataset name, preprocessing details, patient-level split construction, inclusion/exclusion criteria, or provenance of the 5,000 cases, use the underlying run records and dataset documentation. The current three-page summary does not contain those details, so they should not be invented.

This distinction is important because reproducibility in medical imaging depends on exact dataset provenance and split methodology.

19. Questions the Journal Writer May Ask

Question	Answer
What is Med-Nexus?	A medical multimodal research system integrating Project 2 retrieval-aware reliability/conflict analysis with Pro
What is the main research gap?	The system explicitly studies evidence reliability/conflict and task-oriented attention contribution together rath
Why retrieval?	Retrieved evidence can be useful but can also be noisy or contradictory; the project measures behaviour under
Why intervention?	Attention overlap alone does not establish task importance. Intervention measures the effect of masking an int
What is the ΔLoss formula?	$\Delta_{ij} = L(\text{mask}(i,j)) - L(\text{dense})$.
How many seeds?	Five: 42, 123, 456, 789 and 2024.
What GPU?	One NVIDIA A100-SXM4 80GB on Vast.ai.
How large was the held-out evaluation set?	N=5,000 paired image-text clinical cases.
Best documented AUROC?	0.962 ± 0.001 for Full Med-Nexus.
Conflict detection?	97.4% for Full Med-Nexus.
Retrieval?	96.8% Top-1 and 99.6% Top-5 for Full Med-Nexus.
Efficiency?	11.2 ms/sample, 3.1 GB inference VRAM, 46.2% token gating and 89.3 samples/s for Full Med-Nexus.
Latency reduction?	54.8% relative to MN-RUN-05 under the documented comparison.
VRAM reduction?	45.6% relative to MN-RUN-05 under the documented comparison.

Question	Answer
Clinically validated?	Not established by these experiments.

20. Safe Scientific Wording

Use: “The experiments report...”

Use: “Under the stated evaluation protocol...”

Use: “The Full Med-Nexus configuration achieved...”

Use: “The ablation indicates/measures...”

Use: “The results support the evaluated hypothesis under these conditions...”

Avoid: “This proves the model is clinically safe.”

Avoid: “This replaces radiologists.”

Avoid: “This is universally state-of-the-art.”

Avoid: “46.2% of the model was removed.”

21. Five-Minute A-to-Z Call Script

Opening:

“Med-Nexus is the application-level medical multimodal research system integrating Project 2 and Project 4. Project 2 contributes retrieval-aware evidence handling and cross-modal conflict analysis. Project 4 contributes intervention-based attention contribution analysis and contribution-aware long-context gating.”

Problem:

“The motivation is that multimodal medical systems do not only fail because of classification difficulty. Retrieved evidence can be noisy, visual evidence can be ambiguous, and image and text can disagree. We therefore evaluate reliability and conflict explicitly.”

Project 2:

“We stress evidence quality through retrieval-related evaluation, visual ambiguity and cross-modal conflict, and measure retrieval, conflict and calibration behaviour.”

Project 4:

“Instead of treating attention overlap as redundancy, we intervene on selected interactions and measure task-loss change. The documented quantity is $\Delta_{ij} = L(\text{mask}(i,j)) - L(\text{dense})$. This gives a task-oriented contribution signal.”

System:

“The pipeline is medical image and clinical/query text, visual representation, evidence retrieval, reliability/conflict analysis, multimodal fusion, contribution-aware long-context attention, diagnostic prediction and calibration/efficiency evaluation.”

Evaluation:

“We used one A100-SXM4 80GB on Vast.ai, five seeds—42, 123, 456, 789 and 2024—and a

held-out paired image-text clinical evaluation set of 5,000 cases. Maximum sequence length was 2,048, with BF16, batch 32 per GPU, effective batch 64, 100 epochs, 12,500 steps, 1,000 warmup steps, AdamW and peak learning rate 3×10^{-4} ."

Results:

"The main ablation goes from 8B and 12B Baseline s to a 32B Dense baseline, overlap suppression, intervention-based Δ Loss, and finally Full Med-Nexus. The Full configuration has 33.40B parameters and reports AUROC 0.962 ± 0.001 , AUPRC 0.941 ± 0.002 , F1 0.918 ± 0.001 and accuracy 0.938 ± 0.001 ."

Reliability:

"It reports Top-1/Top-5 retrieval of 96.8%/99.6%, conflict detection of 97.4%, ECE 0.011 ± 0.001 and Brier 0.045 ± 0.001 ."

Efficiency:

"It reports 46.2% token gating, 11.2 ms/sample, 3.1 GB inference VRAM and 89.3 samples/s. Relative to MN-RUN-05, the documented latency reduction is 54.8% and inference-VRAM reduction is 45.6%."

Closing:

"The results are research evaluation results under the specified protocol. They do not establish prospective clinical validation, regulatory approval or replacement of clinicians."

22. One-Minute Quick Reference

Project	Med-Nexus
Core mapping	Project 2 + Project 4
Project 2	Retrieval-aware robustness + conflict
Project 4	Intervention contribution + efficient attention
GPU	1× A100-SXM4 80GB, Vast.ai
Seeds	42 / 123 / 456 / 789 / 2024
Held-out cases	5,000 paired image-text clinical cases
Sequence length	2,048
Final model	33.40B parameters
AUROC	0.962 ± 0.001
AUPRC	0.941 ± 0.002
F1	0.918 ± 0.001
Accuracy	0.938 ± 0.001
Retrieval	96.8% / 99.6% Top-1 / Top-5
Conflict detection	97.4%
ECE	0.011 ± 0.001
Brier	0.045 ± 0.001

Project	Med-Nexus
Token gating	46.2%
Latency	11.2 ms/sample
Inference VRAM	3.1 GB
Throughput	89.3 samples/s

23. Glossary

AUROC: Area under the receiver operating characteristic curve; measures ranking/discrimination across thresholds.

AUPRC: Area under the precision-recall curve.

F1: Harmonic mean of precision and recall.

Sensitivity: True-positive rate.

Specificity: True-negative rate.

ECE: Expected Calibration Error; measures discrepancy between confidence and observed correctness.

Brier score: Probabilistic prediction error metric.

RAG / retrieval: Retrieving external or stored evidence to provide additional context to a model.

Cross-modal conflict: Disagreement between information sources such as image-derived and retrieved textual evidence.

Intervention: Deliberately masking/removing a selected interaction and measuring the resulting task effect.

Δ Loss: Difference between intervened/masked loss and dense-reference loss under the documented definition.

Attention overlap: Similarity/overlap between attention patterns; not automatically equivalent to redundancy.

Gating: Suppressing or reducing selected computation according to a gating signal.

Token gating: The documented fraction of tokens subject to the gating mechanism.

Ablation: Controlled comparison in which components are added/removed to study their measured effect.

Held-out test set: Evaluation data kept separate from the training process under the declared protocol.

24. Final Checklist Before the Call

- Keep this PDF open during the call.
- Keep the current GitHub repository open.
- Keep the final results table available.
- Memorize the Project 2 → Project 4 → Med-Nexus relationship.
- Memorize the six MN-RUN configurations.
- Memorize the Full Med-Nexus headline metrics.

- Know the Δ Loss formula and its exact sign interpretation.
- Know the A100 / five-seed / N=5,000 protocol.
- Do not confuse software tests with research validation.
- Do not invent dataset provenance or preprocessing details not present in the source records.
- Do not call the work clinically validated.
- Use “under the stated protocol” when discussing empirical results.

25. Source-of-Truth Note

This A-to-Z guide was prepared from the current Med-Nexus_Complete_Research_Documentation supplied in the project materials, including its final architecture description, research questions, experimental protocol, main ablation, grounding/conflict/calibration results, efficiency results, limitations and handover status.

An older supervisor document exists in the project materials and describes an earlier planning/validation stage. For the final results discussion, the current Complete Research Documentation should be treated as the primary summary source. If a journal writer asks for information not present in that document—especially exact dataset provenance, patient-level splitting, raw run identifiers, or detailed preprocessing—retrieve it from the underlying experiment records rather than filling the gap from memory.

Final principle: Explain what was measured, how it was measured, what changed between controlled configurations, and what the results do—and do not—establish.

END OF A-TO-Z MED-NEXUS HANDOVER GUIDE